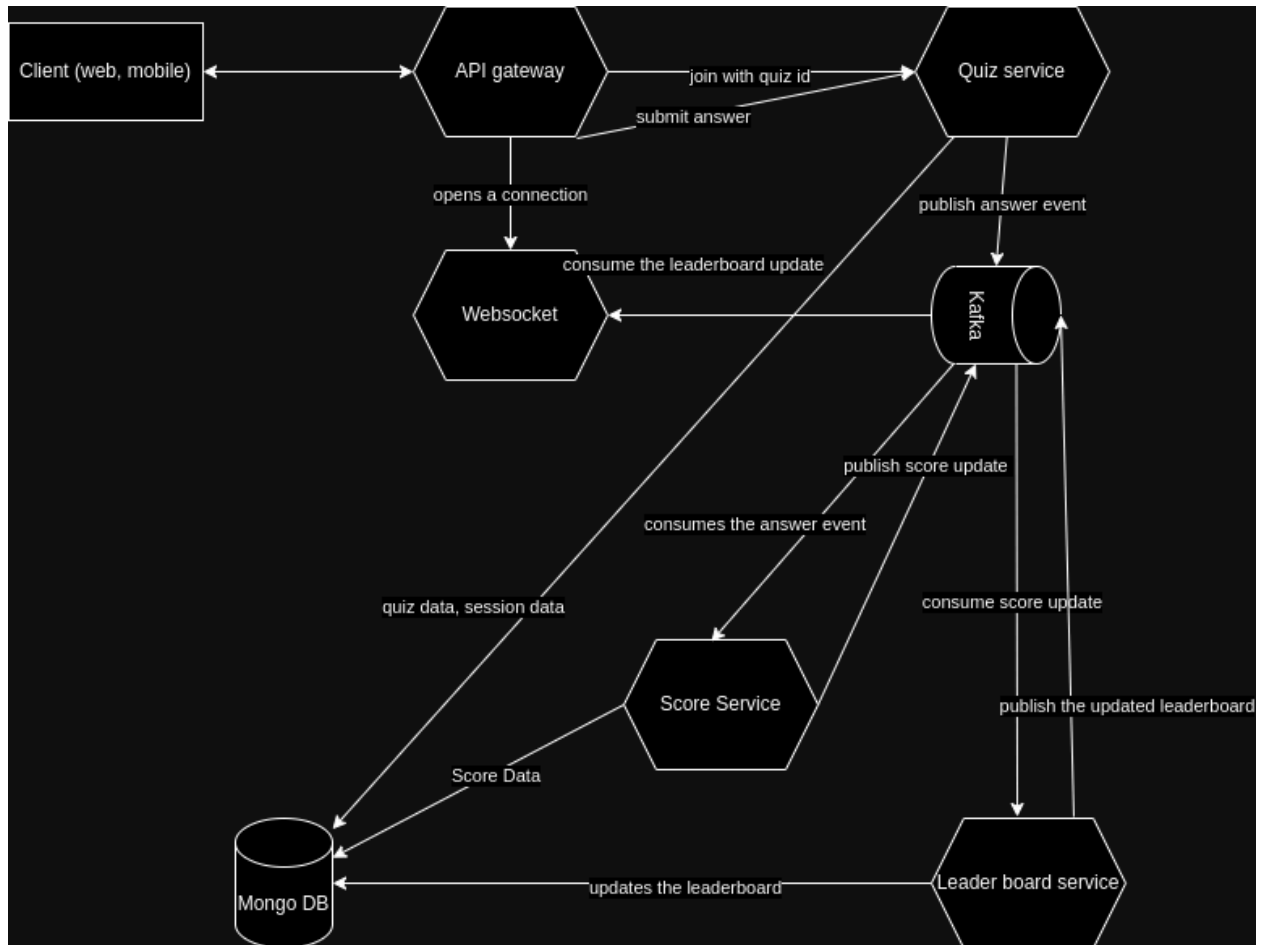


# System Design

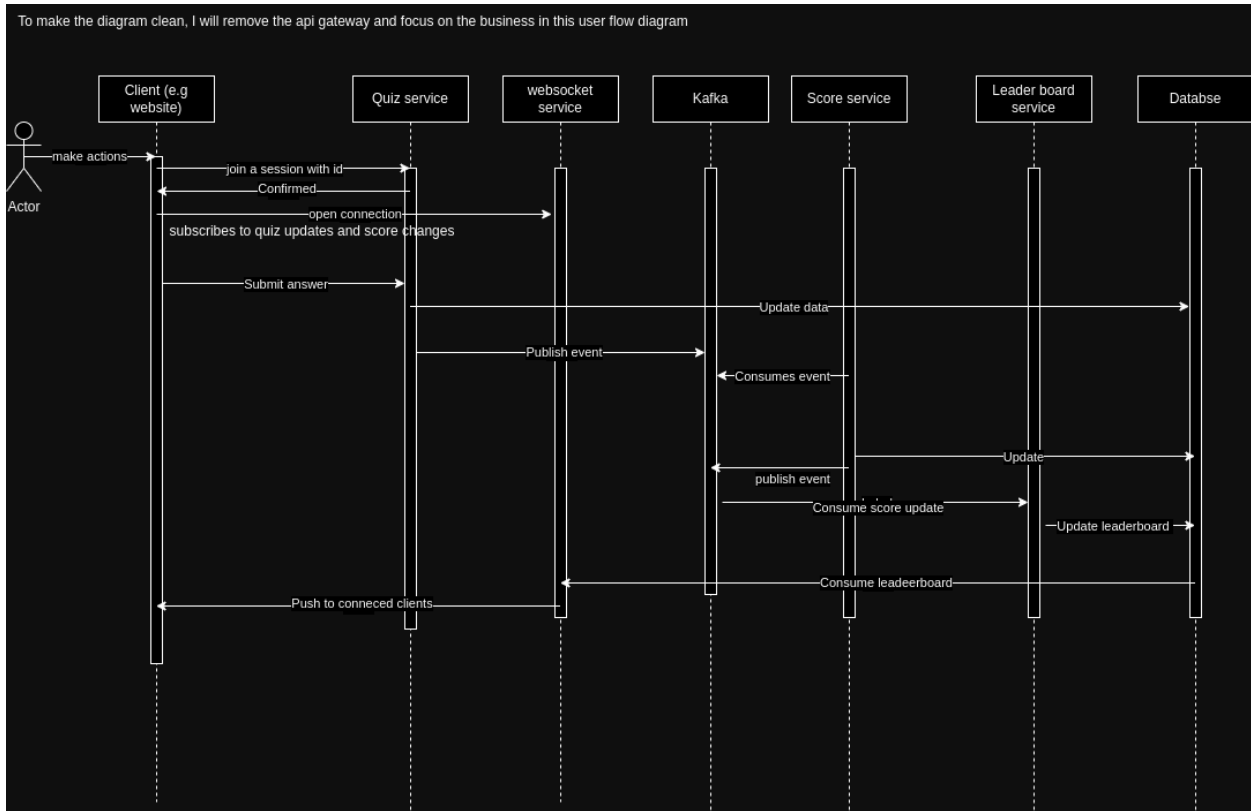
## 1. Architecture Overview Diagram



## 2. Component Description

- **Client Application**: Users can join quizzes, submit answers, and view the leaderboard.
- **Server-side services (Quiz, Leaderboard, and Score)**: Handles business logic, processes quiz answers, updates scores, and manages quiz sessions.
- **Database (MongoDB)**: Persists data related to users, quizzes, and scores.
- **Apache Kafka**: Asynchronous communication between services.
- **WebSocket Service**: Ensures real-time updates of scores and leaderboard.
- **External Services**: Provides additional functionalities like authentication, monitoring, etc.

## 3. Data Flow



#### 1. User Joins Quiz:

- A user sends a request to the Quiz Service to join a quiz using a unique quiz ID.
- The Quiz Service verifies the quiz ID and responds to the client confirming the user's participation.

#### 2. Real-time Quiz Participation:

- The client opens a WebSocket connection to the WebSocket Service.
- The service keeps track of connected users in Redis
- The client subscribes to quiz updates and score changes via this connection.

#### 3. User Submits Answer:

- The client sends the user's answer to the Quiz Service via an HTTP request.
- The Quiz Service publishes the answer event to Kafka.
- Keep track of persistent data in the database

#### 4. Score Update:

- The Score Service consumes the answer event from Kafka.
- The Score Service processes the answer, updates the score in the MongoDB database, and publishes a score update event to Kafka.

#### 5. Leaderboard Update:

- The Leaderboard Service consumes the score update event from Kafka.

- The Leaderboard Service retrieves the latest scores and updates the leaderboard in the MongoDB database.
  - The Leaderboard Service publishes the updated leaderboard event to Kafka.
6. **Real-time Leaderboard:**
- The WebSocket Service consumes the leaderboard update event from Kafka.
  - The WebSocket Service broadcasts the updated leaderboard to all connected clients (get from Redis)

## 4. Technologies and Tools

1. **Service side (Leader board, Score, Quiz, and WebSocket Service): Scala + Akka HTTP**
  - a. **Real-Time Communication:** WebSockets are essential for real-time bidirectional communication, allowing the server to push updates to the client as soon as they happen.
  - b. **Scalability:** Akka HTTP, a highly performant library for building HTTP-based services in Scala, integrates well with Akka Streams, making it a good choice for handling numerous WebSocket connections efficiently.
  - c. **Concurrency Handling:** Akka HTTP leverages the actor model provided by Akka, which is ideal for managing concurrent connections and messages without the complexity of traditional thread-based approaches.
  - d. **Performance:** Scala and Akka are well known for high-performance
  - e. **Reliability:** A strong-type, functional language is very reliable.
  - f. **Last, the ecosystem:** The JVM stack is mature, rich, and stable, especially back-compatibility when upgraded so maintainability is easy and quite cheap.
2. **Client-side technologies**
  - a. Web: any framework (Angular) or library (vuejs, react) will be ok
  - b. Mobile: native android and iOS
3. **Apache Kafka as the backbone for service communication**
  - a. **Decoupling Services**
  - b. **Scalability**
  - c. **Fault Tolerance**
4. **Database (MongoDB):** Persists data related to users, quizzes, and scores with high-performance, easy to scale with heavily loaded.
5. **Caching mechanism:** Redis
6. **Authentication mechanism:** Any 3rd party tool will be fine, e.g Auth0
7. **Monitoring and Observability (Prometheus and Grafana)**
  - a. **Metrics Collection:** Prometheus is a powerful monitoring and alerting toolkit that can collect metrics from your application, including custom metrics from your Scala services.

- b. **Alerting:** Prometheus integrates well with alerting systems, allowing you to set up alerts based on specific thresholds or anomalies in your application's performance.
- c. **Visualization:** Grafana is a popular open-source platform for monitoring and observability that works seamlessly with Prometheus. It provides customizable dashboards to visualize metrics in real time.
- d. **Scalability:** Both Prometheus and Grafana are designed to handle high volumes of metrics and data, making them suitable for applications with significant load.

#### 8. Logging (ELK Stack: Elasticsearch, Logstash, Kibana)

- a. **Centralized Logging:** The ELK stack allows you to collect, process, and visualize logs from various sources in one centralized location.
- b. **Search and Analysis:** Elasticsearch provides powerful full-text search capabilities, making it easy to query and analyze logs.
- c. **Data Processing:** Logstash allows you to ingest, transform, and ship logs from multiple sources to Elasticsearch.
- d. **Visualization:** Kibana offers rich visualization options, enabling you to create dashboards that provide insights into your application's logs.
- e. **Scalability:** The ELK stack can handle large volumes of log data, making it suitable for applications with extensive logging requirements.

#### 9. Containerization and Orchestration (Docker and Kubernetes)

- a. **Consistency:** Docker ensures that your application runs consistently across different environments by packaging it with all its dependencies.
- b. **Isolation:** Docker containers provide process isolation, improving security and resource management.
- c. **Scalability:** Kubernetes automates the deployment, scaling, and management of containerized applications. It helps ensure the high availability and resilience of your services.
- d. **Orchestration:** Kubernetes provides advanced orchestration features, such as automated rollouts and rollbacks, service discovery, and load balancing, which are essential for managing microservices in production.